

Cache Attacks and Countermeasures – the Case of AES

Elena Frank, Frederik Safner, Lukas Kaibel

July 5, 2023

Contents

1	Introduction	2
2	Synchronous Known-Data Attacks	2
3	Asynchronous Attacks	3
4	Countermeasures	3
5	Conclusion	4

1 Introduction

The paper discusses cache attacks as a type of side-channel attacks that leverage cross-process information leaking through CPU caches.

The attacks are possible due to how *set-associative memory caches* are implemented. Their structure limits the mapping of memory addresses to the cache in the sense a) always full memory blocks are cached, and b) that a memory block can only be cached in specific cache lines.

The attacked cipher is the AES cipher, where the attacks leverages how it is implemented in performance-oriented software. Concrete, the implementation pre-computes Lookup-tables. These tables are then used to realize the algebraic Shift-Rows, Mix-Columns and Sub-Bytes operations needed in each round of the AES computation.

The paper describes multiple attacks where a process can obtain information about another process' memory access patterns by leveraging said implementation details of the the memory caches and AES implementation. Finally, the paper also discusses countermeasures for such attacks.

2 Synchronous Known-Data Attacks

Synchronous attacks are applicable if the plaintext or ciphertext is known. The main assumption is that the attacker operates synchronously with the encryption on the same processor. Therefore, the attacker has to know when some encryption starts. This is a realistic assumption, as you can see in the following example. A virtual device provides encrypted storage to the physical device, having a normal filesystem mounted on top. A user being able to write on any file might now trigger an encryption on a chosen plaintext. Using the presented attack the attacker is able to recover the key used for encryption.

One-Round Attack: The simplest synchronous attack exploits the fact, that given knowledge of a plaintext byte, information about the accessed index exposes information about the key at this position. The main idea basically is the following. The attacker collects measurements related to the encryption. These contain information whether specific memory blocks are accessed during the encryption process. By comparing the samples with ideal predictions, the attacker can filter out noise and may also make educated guesses about parts of the encryption key. To estimate the likelihood of a specific key being correct, the candidate score gets introduced. It is higher when the measurement scores align with the expected behavior of memory block accesses. However, even though the One-Round Attack demonstrates its effectiveness in revealing the top four bits, it is limited by the need for statistical analysis and the reliance on specific conditions and measurement.

Two-Rounds Attack: In comparison two the one-round attack, the two round attack analyzes also the second round of AES encryption. This especially includes the non-linear mixing of the Rijndael cipher to narrow down the possible values of key bytes. An ideal predicate is used as a condition that represents the expected behavior in an ideal scenario. With partial knowledge of the key bytes from round one, the unknown low bits of certain key bytes in the second round affect only the low bits of a specific intermediate value. By enumerating candidate values, the correct candidate can be identified.

Measurement via Evict+Time: In general for Evict and Time, the attacker evicts memory from the cache using an evictions set. If the victim accesses the shared memory, it overwrites the cache. Finally the attacker accesses the shared memory. If the access is quick, it has been accessed in between by the victim. Otherwise it has not. In the case of AES encryption the timing differences are observed during the encryption process. This reveals information about specific table indices being accessed or not. The information is later used for the extraction of measurement scores, which is essential to the recovery of the encryption key.

Measurement via Prime+Probe: Prime and Probe is also a common building block for side channel attacks. The attacker flushes monitored cache lines (using `clflush` on x86) and waits for the victim to execute its code. The attacker then measures the access time to some cache lines, to decide whether the victim has accessed the cache lines. The timing is done using clock cycle counter `rdtsc`. Here the cache is also primed by reading values from all memory blocks. After triggering an encryption of a chosen plaintext the cache is probed to check for cache misses in different cache sets. This way the attacker can infer which memory blocks have been accessed during encryption. In comparison to Evict+Time, the method is less error-prone as it is less sensitive to timing variations.

The One-Round Round Attack, the Two-Round Attack, Evict+Time and Prime+Probe can be used together, as part of a bigger attack. Combined, they further refine the key byte candidates, which provides a more effective and efficient key recovery.

3 Asynchronous Attacks

Synchronous attacks require that the attacker can interact with the encryption code. For asynchronous attack this is not necessary: a process doesn't have to interact with encryption program and can still obtain information simply by running on the same CPU.

It is assumed that multithreading is supported and that the attacker process runs in parallel with the victim. The attacker can analyze the frequency score of the victims memory accesses, i.e. how often it accesses a memory location. The preliminary is that the attacker's and the victim's used memory blocks map to the same cache lines. The attacker measures the time it takes to retrieve memory blocks when it runs alone, vs when it runs in parallel with the victim. Depending on how often the victim access its memory blocks, the attacker's process gets cache misses when loading its blocks, which shows through an increased access time. By measuring over a larger number of access, the frequency score for the victim can be obtained.

This frequency score for each AES Lookup Table can be used together with the frequency information about the encrypted text (e.g. that if it's an english text, it is very likely that a bytes has there highest nibble set to 6, which represent a letter between 'a' and 'p') to obtain information about the highest nibble in some key bytes. The paper described that through this method, 45.66 bits of the secret key could be obtained.

It also proposes that if additionally the order of accesses could be analyzed by the attacker, even more bits can be obtained. Furthermore, it describes an way to remove the limitation that the host system must be multithreaded: by exploiting interrupt mechanism an attacker can pause the encryption process and analyze the state of the cache (again through loading memory blocks and measuring cache misses) again as discussed. For processors with multiple layers of caching i.e. where L1 caches are private to a core, the paper claims that the attacks may still be possible if the cryptographic code to L2 / L3 caches that are shared among cores.

4 Countermeasures

Several methods can be considered to mitigate information leakage in cryptographic systems, each with unique trade-offs:

Avoiding Table Lookups: Techniques include using logical operations instead of table lookups, using bitslice implementations to execute multiple encryptions simultaneously through vectorization, or placing lookup tables in registers instead of cache. These methods may reduce performance or be architecture-dependent.

Alternative Lookup Tables: By using smaller or combined lookup tables, the chance of a given memory block not being accessed during encryption is reduced. This method is more resistant to synchronous attacks but still vulnerable to asynchronous attacks, and is most effective when used with other countermeasures.

Data-Oblivious Memory Access Pattern: This involves altering memory access patterns to be independent of data. Techniques include reading all entries from relevant tables, memory shuffling or permutation, or adding spurious memory accesses. All these methods result in some performance slowdown.

Algorithmic Masking Techniques: These randomize data-dependent operations. Current AES masking methods are either insecure or significantly slower.

Cache State Normalization and Process Blocking: This method involves normalizing the cache state before and after encryption and is effective against synchronous attacks but not asynchronous ones.

Disabling Cache Sharing or Caching Mechanism: Disabling cache sharing can prevent cache state effects from crossing process boundaries but may require extensive resources. A more drastic measure is disabling the CPU's caching mechanism altogether, which significantly affects performance. A more feasible alternative is the "no-fill" mode, which requires privileged mode.

Dynamic Table Storage: This method involves randomizing table lookup locations or changing table order multiple times during encryption. The effectiveness of this method depends on the architecture.

Hiding Timing Information: This involves adding random delays or normalizing all timings to a fixed value, but it doesn't provide complete protection against certain types of attacks.

Selective Round Protection: This method focuses on protecting specific rounds of encryption while leaving the internal rounds unprotected.

OS-Level Support: Some countermeasures may require OS-level support or privileged operations. One solution could be to provide secure execution of cryptographic primitives as OS services. An alternative is providing a secure execution facility to user processes.

These methods represent various trade-offs between security and performance. Due to the complexity and architecture-dependency, secure implementation requires a delicate balance and no one-size-fits-all solution exists.

5 Conclusion

The study examined the vulnerability of cryptographic systems like AES to cache state analysis attacks, noting that methods using data-dependent memory access or large lookup tables are particularly at risk. System-level concerns were also raised as both multi-user systems and single-user machines could be affected. Some countermeasures were proposed but none were entirely efficient, leaving the mitigation of such attacks an open problem.